

Tone: Purpose and Value

Learn your voice once. Sound like yourself everywhere.

The problem Tone exists to solve

Marketing teams already use AI to write — and mostly hate the results. Generic AI tools have two structural flaws:

1. **They sound like AI.** The same over-polished, buzzword-heavy register every company on the internet is now publishing. For a brand, sounding generic is worse than being slow: it erodes the one asset content is supposed to build — a recognizable voice.
2. **They forget.** Every session starts from zero. Someone pastes the style guide into the prompt again, explains the audience again, lists the banned words again. Multiply that across a team and a content calendar, and “AI speed” quietly disappears into re-prompting and heavy rewriting.

What Tone is

Tone is an AI writing workspace whose core is not the generator — it is the **Brand Voice Profile**: a persistent, structured, editable model of how your company actually sounds.

A profile is built once, in a five-minute onboarding: 3–5 real writing samples, your website, and a short quiz covering tone sliders (formal–casual, careful–bold, concise–detailed), words you love, words you ban, your audience, and the competitors you refuse to sound like. Tone distills that into named voice traits you can inspect and correct.

From then on, **every single generation is conditioned on the active profile**. Nobody re-explains the brand, ever. And the profile is not static: thumbs-up/down feedback and “more like this” edits feed back into it, and a visible **Voice Accuracy score** climbs as it learns. The tool gets more yours every week you use it.

Who it is for

Tone is designed for **B2B SaaS and tech marketing teams of 1–20 people** — content leads, founders who post, demand-gen marketers, and support teams — though the vertical is a single configurable setting, so any niche fits.

These teams share a profile: they publish across many surfaces (blog, LinkedIn, email, landing pages), they care intensely about credibility with a sophisticated audience, and they have no time for a rewriting bottleneck.

How it helps, concretely

- **Consistency across every surface.** One workspace holds multiple voices — “Company blog”, “Founder LinkedIn”, “Lifecycle email” — each trained separately, so the founder does not sound like the docs and the docs do not sound like ChatGPT.
- **First drafts that start at 80%, not 40%.** Because generation is conditioned on real samples, banned words, and audience, the heavy rewrite pass mostly disappears. Three distinct takes arrive side by side; you pick, tweak, ship.
- **Editing at the speed of intent.** Shorten, Expand, Regenerate, and More on-brand are one-click, voice-aware operations, with live tone sliders that re-shape output in place.
- **A content pipeline, not a chat log.** Drafts save with their brief, format, and voice attached, and move through Draft → Approved → Published with team roles (owner, editor, viewer) enforcing who can do what.
- **Compounding value.** Every piece of feedback makes the profile — the team’s asset — smarter. That is also why teams stay: leaving Tone means abandoning a trained voice that no fresh tool can replicate on day one.

The bottom line

Generic AI makes everyone sound the same. Tone makes AI sound like *you* — and keeps getting better at it. Teach it your voice once, and every blog post, LinkedIn post, email, and landing page after that ships faster **and** sounds unmistakably like your brand.